



<http://algs4.cs.princeton.edu>

4.4 SHORTEST PATHS


- *APIs*
- *shortest-paths properties*
- *Dijkstra's algorithm*
- *edge-weighted DAGs*
- *negative weights*

Shortest paths in an edge-weighted digraph

Given an edge-weighted digraph, find the shortest path from s to t .

edge-weighted digraph

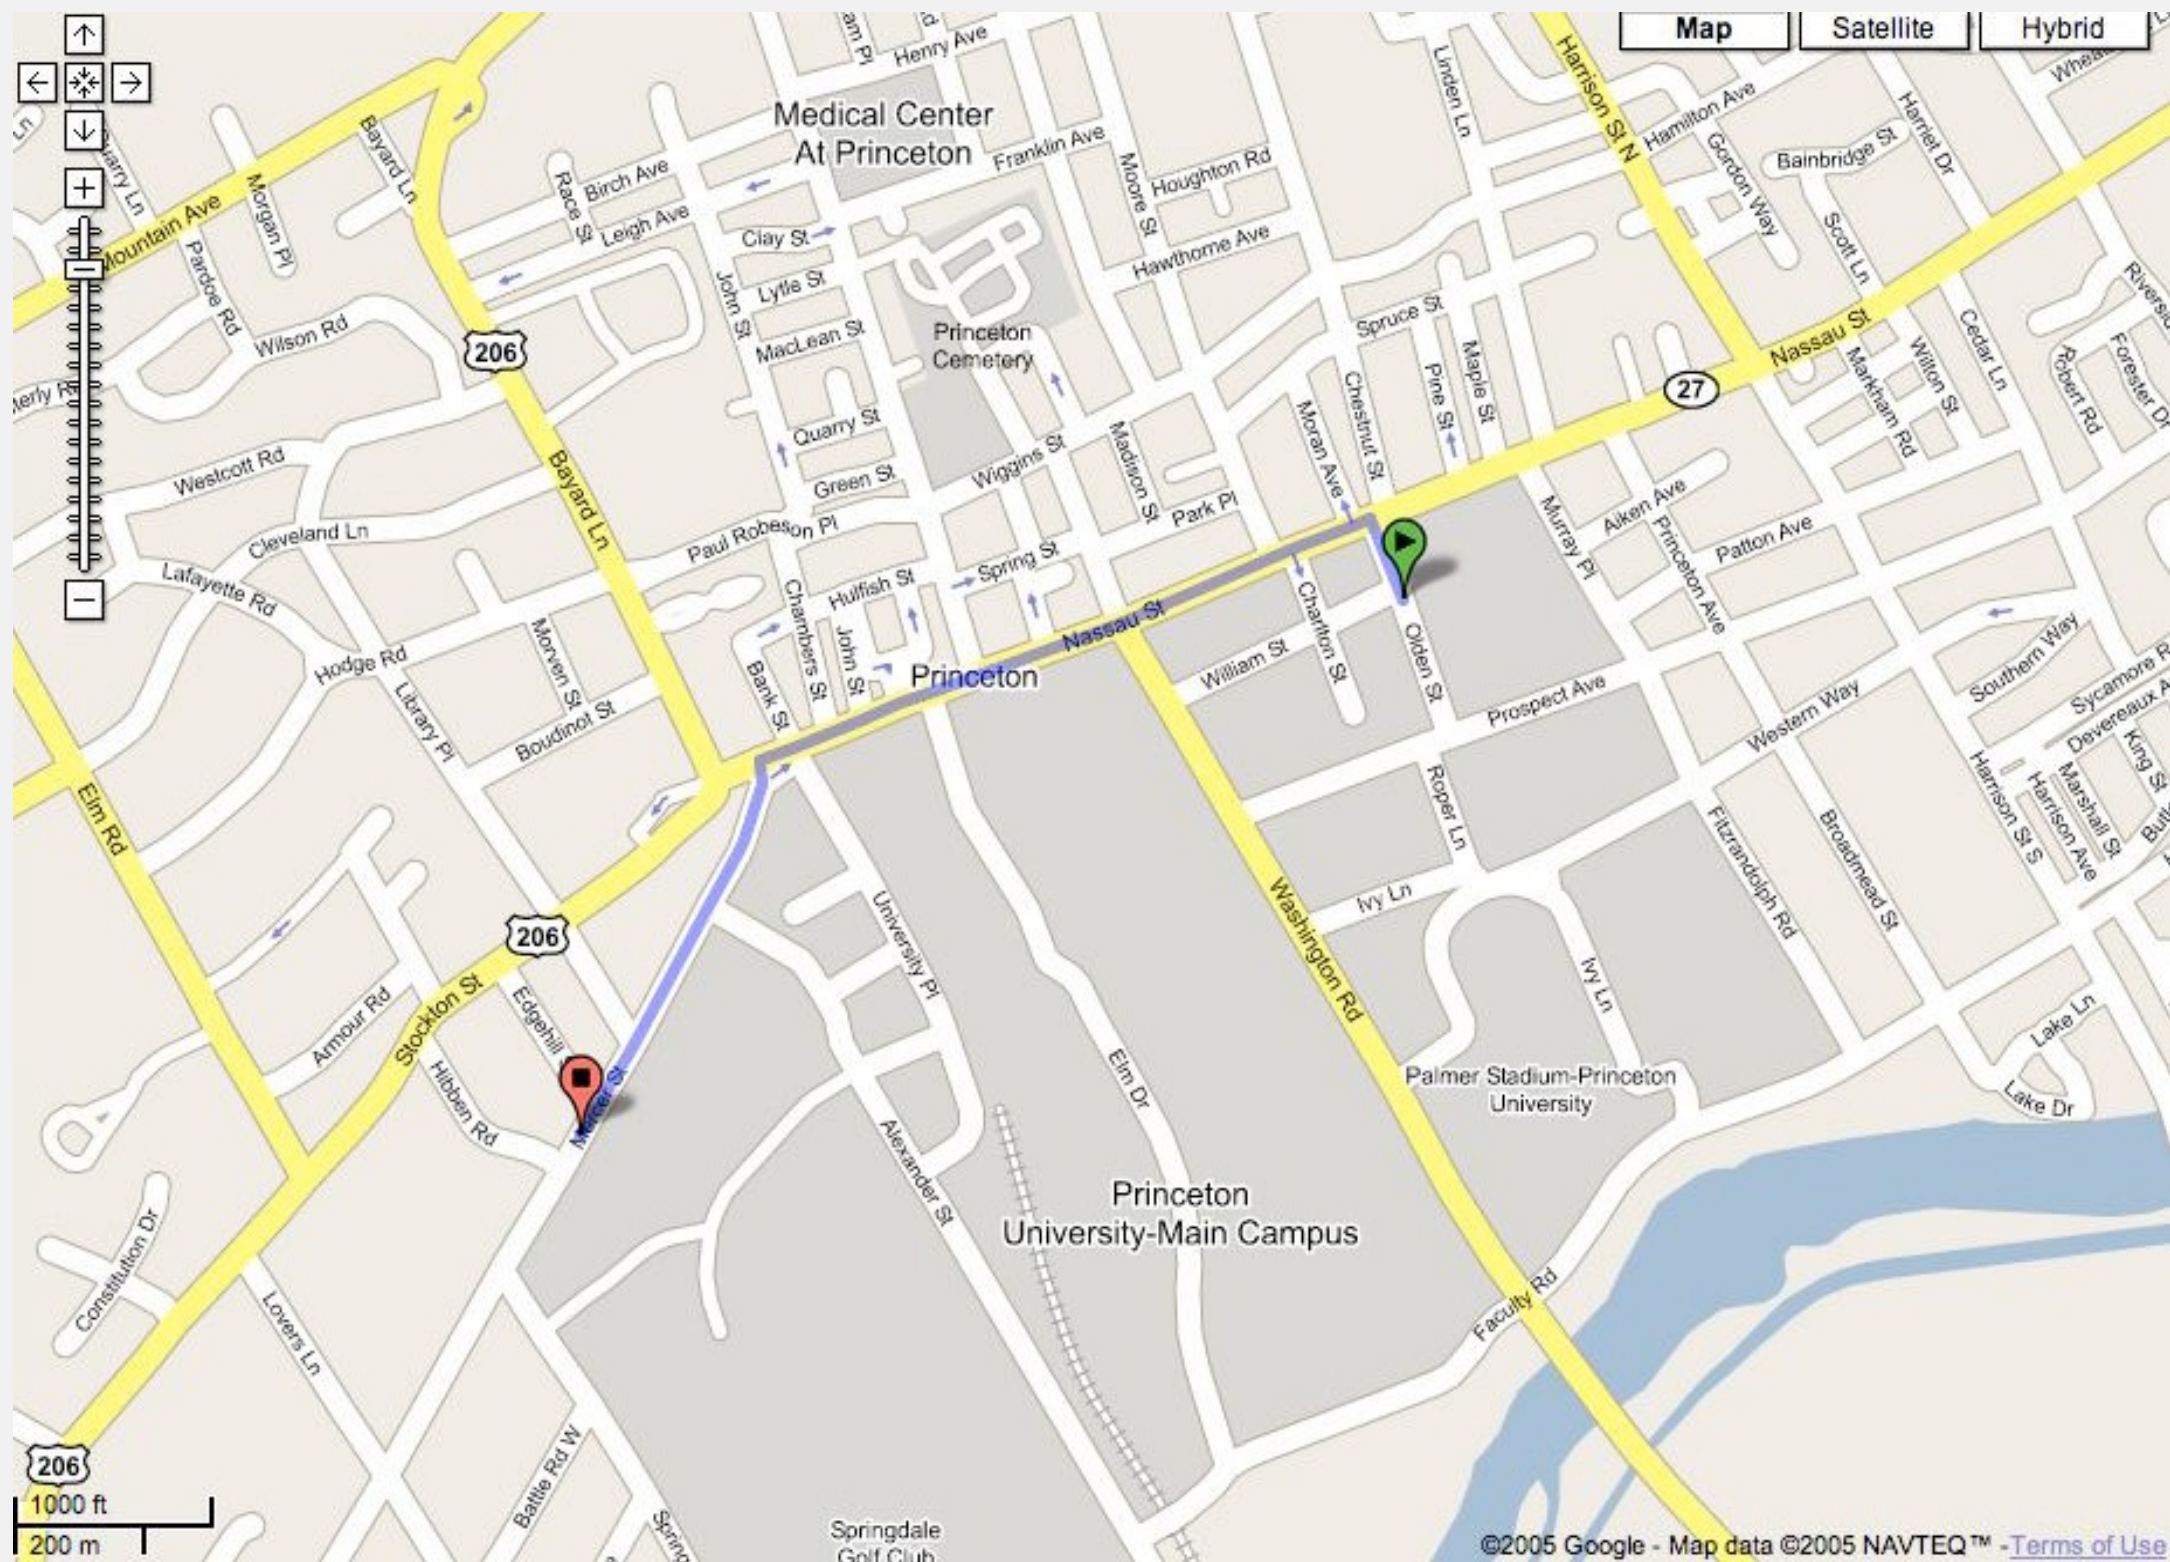
4→5	0.35
5→4	0.35
4→7	0.37
5→7	0.28
7→5	0.28
5→1	0.32
0→4	0.38
0→2	0.26
7→3	0.39
1→3	0.29
2→7	0.34
6→2	0.40
3→6	0.52
6→0	0.58
6→4	0.93



shortest path from 0 to 6

0→2	0.26
2→7	0.34
7→3	0.39
3→6	0.52

Google maps



Shortest path applications

- PERT/CPM.
- Map routing.
- **Seam carving**.
- Texture mapping.
- Robot navigation.
- Typesetting in TeX.
- Urban traffic planning.
- Optimal pipelining of VLSI chip.
- Telemarketer operator scheduling.
- Routing of telecommunications messages.
- Network routing protocols (OSPF, BGP, RIP).
- Exploiting **arbitrage** opportunities in currency exchange.
- Optimal truck routing through given traffic congestion pattern.



http://en.wikipedia.org/wiki/Seam_carving



Reference: Network Flows: Theory, Algorithms, and Applications, R. K. Ahuja, T. L. Magnanti, and J. B. Orlin, Prentice Hall, 1993.

Shortest path variantat

Cilat kulme?

- **Single source:** prej një kulmi s tek çdo kulm tjetër.
- **Single sink:** prej çdo kulmi tek një kulm t .
- **Source-sink:** prej një kulmi s tek tjetri t .
- **All pairs:** Në mes të gjitha çifteve të kulmeve.

Restrikimet për peshën e segmenteve?

- Peshat jonegative.
- Peshat Euclidiane.
- Peshat arbitrare.

Ciklet?

- Pa directed cikle.
- Pa "cikle negative."



which variant?

Simplifying assumption. Shortest paths prej s tek çdo kulm v ekzistojnë.



<http://algs4.cs.princeton.edu>


4.4 SHORTEST PATHS

- ▶ *APIs*
- ▶ *shortest-paths properties*
- ▶ *Dijkstra's algorithm*
- ▶ *edge-weighted DAGs*
- ▶ *negative weights*

Weighted directed edge API

```
public class DirectedEdge
```

<pre> int v, int w, double weight)</pre>	<i>weighted edge $v \rightarrow w$</i>
<pre> DirectedEdge(int from(),</pre>	<i>vertex v</i>
<pre> int to())</pre>	<i>vertex w</i>
<pre> double weight()</pre>	<i>weight of this edge</i>
<pre> String toString()</pre>	<i>string representation</i>



Idiom for processing an edge e : `int v = e.from(), w = e.to();`

Weighted directed edge: implementation in Java

Similar to Edge for undirected graphs, but a bit simpler.

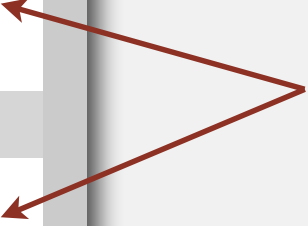
```
public class DirectedEdge
{
    private final int v, w;
    private final double weight;

    public DirectedEdge(int v, int w, double weight)
    {
        this.v = v;
        this.w = w;
        this.weight = weight;
    }

    public int from()
    { return v; }

    public int to()
    { return w; }

    public int weight()
    { return weight; }
}
```



from() and to() replace
either() and other()

Edge-weighted digraph API

```
public class EdgeWeightedDigraph
```

```
    EdgeWeightedDigraph(int V) edge-weighted digraph with V vertices
```

```
    EdgeWeightedDigraph(InputStream in) edge-weighted digraph from input stream
```

```
    void addEdge(DirectedEdge e) add weighted directed edge e
```

```
    Iterable<DirectedEdge> adj(int v) edges pointing from v
```

```
    int V() number of vertices
```

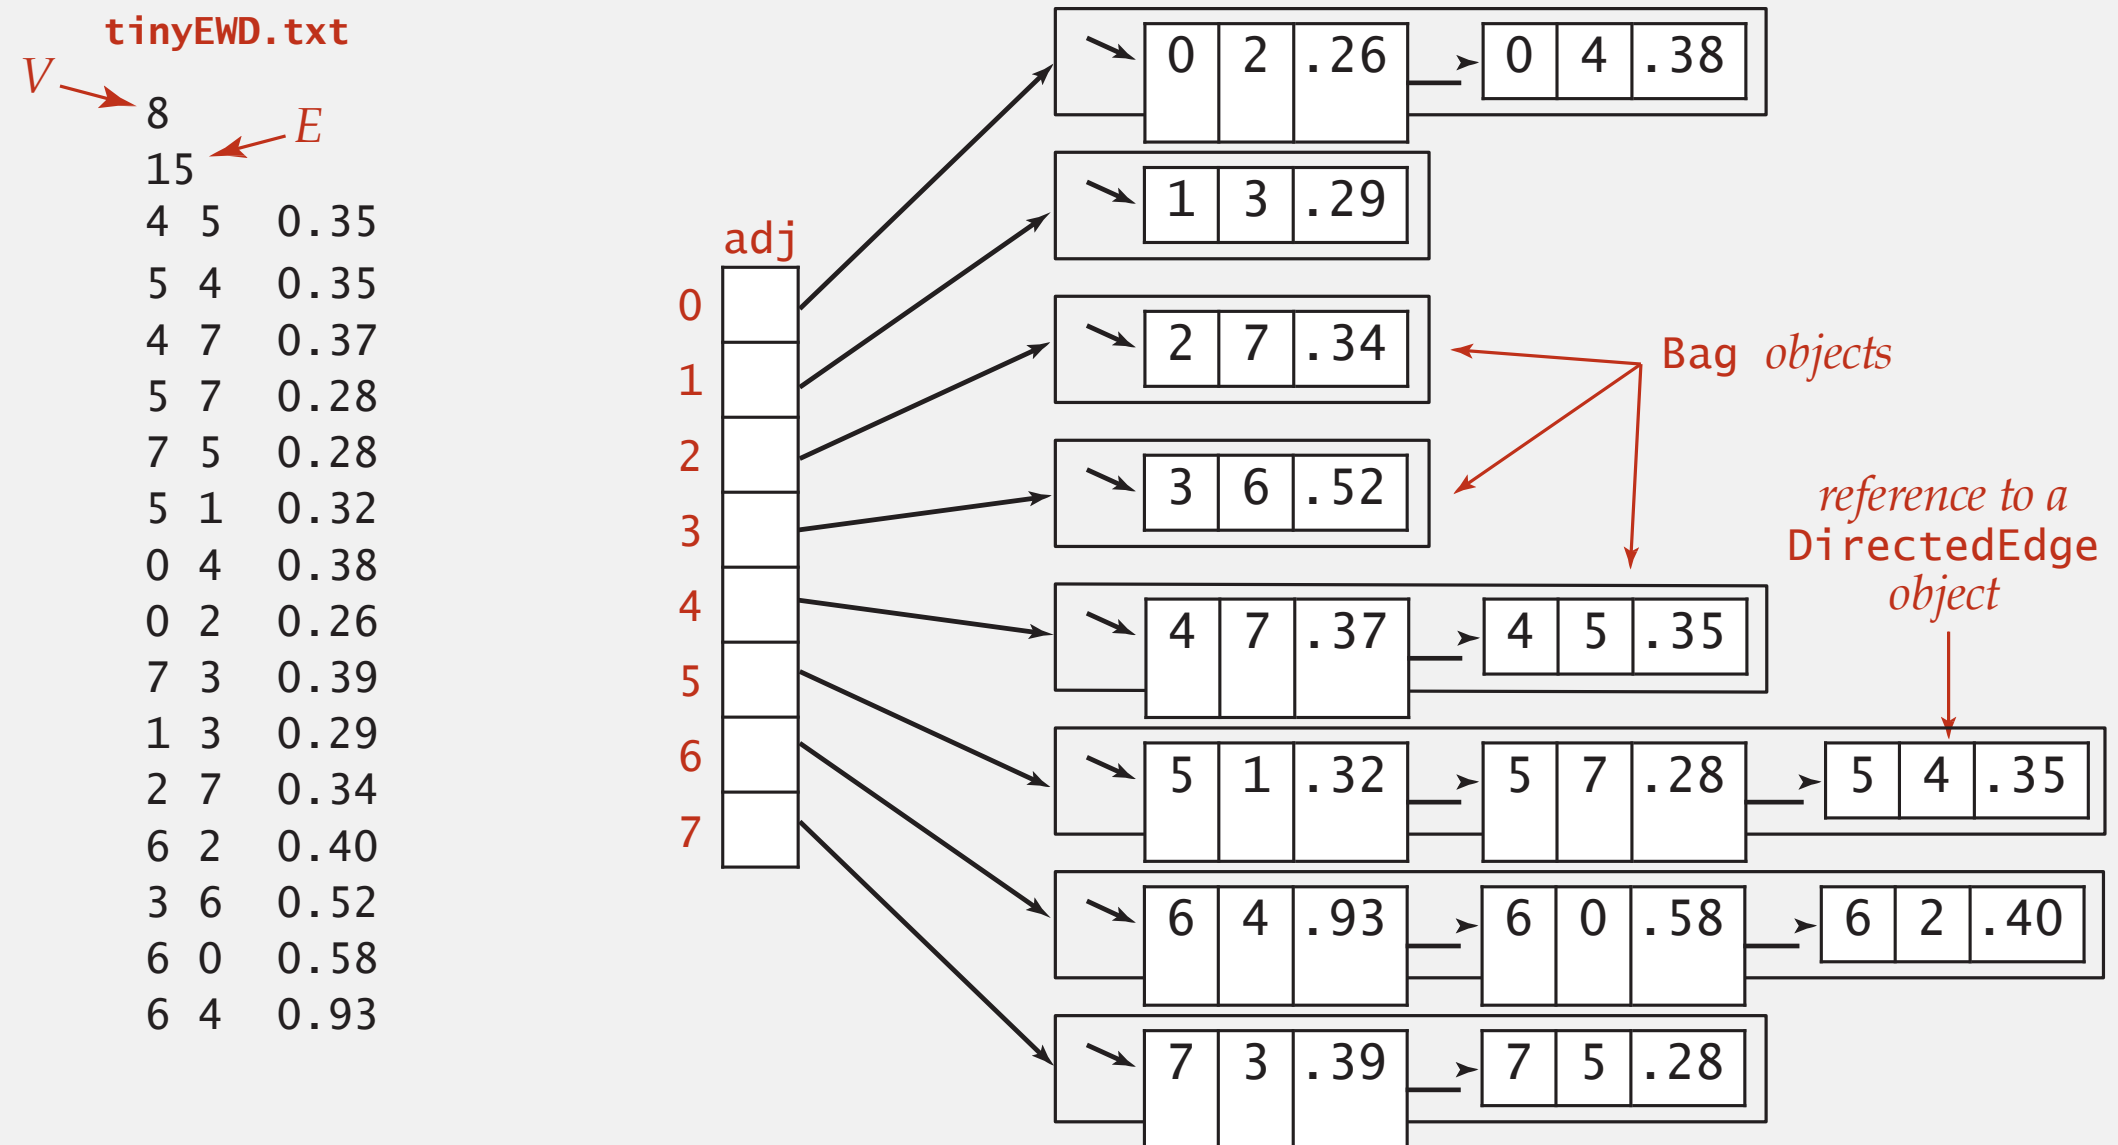
```
    int E() number of edges
```

```
    Iterable<DirectedEdge> edges() all edges
```

```
    String toString() string representation
```

Conventions. Allow self-loops and parallel edges.

Edge-weighted digraph: adjacency-lists representation



Edge-weighted digraph: adjacency-lists implementation in Java

Same as EdgeWeightedGraph except replace Graph with Digraph.


```
public class EdgeWeightedDigraph
{
    private final int V;
    private final Bag<DirectedEdge>[] adj;

    public EdgeWeightedDigraph(int V)
    {
        this.V = V;
        adj = (Bag<DirectedEdge>[]) new Bag[V];
        for (int v = 0; v < V; v++)
            adj[v] = new Bag<DirectedEdge>();
    }

    public void addEdge(DirectedEdge e)
    {
        int v = e.from();
        adj[v].add(e);
    }

    public Iterable<DirectedEdge> adj(int v)
    { return adj[v]; }
}
```

add edge $e = v \rightarrow w$ to
only v 's adjacency list



Single-source shortest paths API

Goal. Find the shortest path from s to every other vertex.

```
public class SP
{
    double SP(EdgeWeightedDigrap G, int s) shortest paths from s in graph G
    double distTo(int v) length of shortest path from s to v
    Iterable <DirectedEdge> pathTo(int v) shortest path from s to v
    boolean hasPathTo(int v) is there a path from s to v?
}
```

```
SP sp = new SP(G, s);
for (int v = 0; v < G.V(); v++)
{
    StdOut.printf("%d to %d (%.2f): ", s, v, sp.distTo(v));
    for (DirectedEdge e : sp.pathTo(v))
        StdOut.print(e + " ");
    StdOut.println();
}
```

Single-source shortest paths API

Goal. Find the shortest path from s to every other vertex.

```
public class SP
{
    double SP(EdgeWeightedDigrap G, int s) shortest paths from s in graph G
    double distTo(int v) length of shortest path from s to v
    Iterable <DirectedEdge> pathTo(int v) shortest path from s to v
    boolean hasPathTo(int v) is there a path from s to v?
}
```

```
% java SP tinyEWD.txt 0
0 to 0 (0.00):
0 to 1 (1.05): 0->4 0.38 4->5 0.35 5->1 0.32
0 to 2 (0.26): 0->2 0.26
0 to 3 (0.99): 0->2 0.26 2->7 0.34 7->3 0.39
0 to 4 (0.38): 0->4 0.38
0 to 5 (0.73): 0->4 0.38 4->5 0.35
0 to 6 (1.51): 0->2 0.26 2->7 0.34 7->3 0.39 3->6 0.52
0 to 7 (0.60): 0->2 0.26 2->7 0.34
```



<http://algs4.cs.princeton.edu>

4.4 SHORTEST PATHS

▸ *APIs*

▸ *shortest-paths
properties*

▸ *Dijkstra's algorithm*

▸ *edge-weighted DAGs*

▸ *negative weights*


Strukturat e të dhënave për single-source shortest paths

Qëllimi. Gjej shortest path prej s tek çdo kulm tjetër.

Vërejtje. **shortest-paths tree** (SPT) ekziston. Pse?

Përfundim. SPT mund të paraqitet me anë të dy arrays me kulme të indeksuara:

- $\text{distTo}[v]$ është gjatësia e shortest path from s tek v .
- $\text{edgeTo}[v]$ është segmenti i fundit në shortest path prej s tek v .



shortest-paths tree from 0

	edgeTo[]	distTo[]
0	null	0
1	5->1 0.32	1.05
2	0->2 0.26	0.26
3	7->3 0.37	0.97
4	0->4 0.38	0.38
5	4->5 0.35	0.73
6	3->6 0.52	1.49
7	2->7 0.34	0.60

parent-link representation

Data structures for single-source shortest paths

Qëllimi. Gjej shortest path prej s tek çdo kulm tjetër.

Observation. **shortest-paths tree** (SPT) ekziston. Pse?

Përfundim. SPT mund të paraqitet me anë të dy arrays me kulme të indeksuara:

- $\text{distTo}[v]$ është gjatësia e shortest path from s tek v .
- $\text{edgeTo}[v]$ është segmenti i fundit në shortest path prej s tek v .

```
public double distTo(int v)
{ return distTo[v]; }


public Iterable<DirectedEdge> pathTo(int v)
{
    Stack<DirectedEdge> path = new Stack<DirectedEdge>();
    for (DirectedEdge e = edgeTo[v]; e != null; e = edgeTo[e.from()])
        path.push(e);
    return path;
}
```

Lehtäsimi i segmentit

Lehtäsimi i segmentit $e = v \rightarrow w$.

- $\text{distTo}[v]$ is length of shortest **known** path from s to v .
- $\text{distTo}[w]$ is length of shortest **known** path from s to w .
- $\text{edgeTo}[w]$ is last edge on shortest **known** path from s to w .
- If $e = v \rightarrow w$ gives shorter path to w through v , update both $\text{distTo}[w]$ and $\text{edgeTo}[w]$.

$v \rightarrow w$ successfully relaxes



Edge relaxation

Relax edge $e = v \rightarrow w$.

- `distTo[v]` is length of shortest **known** path from s to v .
- `distTo[w]` is length of shortest **known** path from s to w .
- `edgeTo[w]` is last edge on shortest **known** path from s to w .
- If $e = v \rightarrow w$ gives shorter path to w through v , update both `distTo[w]` and `edgeTo[w]`.

```
private void relax(DirectedEdge e)
{
    int v = e.from(), w = e.to();
    if (distTo[w] > distTo[v] + e.weight())
    {
        distTo[w] = distTo[v] + e.weight();
        edgeTo[w] = e;
    }
}
```

Generic shortest-paths algorithm

Generic algorithm (to compute SPT from s)

Initialize $\text{distTo}[s] = 0$ and $\text{distTo}[v] = \infty$ for all other vertices.

Repeat until optimality conditions are satisfied:

- Relax any edge.
-

Efficient implementations. How to choose which edge to relax?

Ex 1. Dijkstra's algorithm (nonnegative weights).

Ex 2. Topological sort algorithm (no directed cycles).

Ex 3. Bellman-Ford algorithm (no negative cycles).



<http://algs4.cs.princeton.edu>

4.4 SHORTEST PATHS

▸ *APIs*

▸ *shortest-paths
properties*

▸ *Dijkstra's algorithm*

▸ *edge-weighted DAGs*

▸ *negative weights*

Edsger W. Dijkstra: select quotes


“ Do only what only you can do. ”

“ In their capacity as a tool, computers will be but a ripple on the surface of our culture. In their capacity as intellectual challenge, they are without precedent in the cultural history of mankind. ”

“ The use of COBOL cripples the mind; its teaching should, therefore, be regarded as a criminal offence. ”

“ It is practically impossible to teach good programming to students that have had a prior exposure to BASIC: as potential programmers they are mentally mutilated beyond hope of regeneration. ”

“ APL is a mistake, carried through to perfection. It is the language of the future for the programming techniques of the past: it creates a new generation of coding bums. ”




Edsger W. Dijkstra
Turing award 1972

Edsger W. Dijkstra: select quotes



Dijkstra's algorithm demo

- Le të jenë kulmet duke u rritur për nga distanca prej s (non-tree vertex with the lowest $\text{distTo}[]$ value).
- Shto kulmin në dru dhe lehtëso të gjitha segmentet që dalin nga ai kulm.




an edge-weighted digraph

$0 \rightarrow 1$	5.0
$0 \rightarrow 4$	9.0
$0 \rightarrow 7$	8.0
$1 \rightarrow 2$	12.0
$1 \rightarrow 3$	15.0
$1 \rightarrow 7$	4.0
$2 \rightarrow 3$	3.0
$2 \rightarrow 6$	11.0
$3 \rightarrow 6$	9.0
$4 \rightarrow 5$	4.0
$4 \rightarrow 6$	20.0
$4 \rightarrow 7$	5.0
$5 \rightarrow 2$	1.0
$5 \rightarrow 6$	13.0
$7 \rightarrow 5$	6.0
$7 \rightarrow 2$	7.0

Dijkstra's algorithm demo


- Le të jenë kulmet duke u rritur për nga distanca prej s (non-tree vertex with the lowest $\text{distTo}[]$ value).
- Shto kulmin në dru dhe lehtëso të gjitha segmentet që dalin nga ai kulm.




v	$\text{distTo}[]$	$\text{edgeTo}[]$
0	0.0	-
1	5.0	0→1
2	14.0	5→2
3	17.0	2→3
4	9.0	0→4
5	13.0	4→5
6	25.0	2→6
7	8.0	0→7

shortest t-paths tree from vertex s

Dijkstra's algorithm visualization



Dijkstra's algorithm visualization



Dijkstra's algorithm: Java implementation

```
public class DijkstraSP
{
    private DirectedEdge[] edgeTo;
    private double[] distTo;
    private IndexMinPQ<Double> pq;

    public DijkstraSP(EdgeWeightedDigraph G, int s)
    {
        edgeTo = new DirectedEdge[G.V()];
        distTo = new double[G.V()];
        pq = new IndexMinPQ<Double>(G.V());

        for (int v = 0; v < G.V(); v++)
            distTo[v] = Double.POSITIVE_INFINITY;
        distTo[s] = 0.0;

        pq.insert(s, 0.0);
        while (!pq.isEmpty())
        {
            int v = pq.delMin();
            for (DirectedEdge e : G.adj(v))
                relax(e);
        }
    }
}
```

← relax vertices in order
of distance from s

Dijkstra's algorithm: Java implementation

```
private void relax(DirectedEdge e)
{
    int v = e.from(), w = e.to();
    if (distTo[w] > distTo[v] + e.weight())
    {
        distTo[w] = distTo[v] + e.weight();
        edgeTo[w] = e;
        if (pq.contains(w)) pq.decreaseKey(w, distTo[w]);
        else                  pq.insert      (w, distTo[w]);
    }
}
```

← update PQ

Dijkstra's algorithm: which priority queue?

Depends on PQ implementation: V insert, V delete-min, E decrease-key.

PQ implementation	insert	delete-min	decrease-key	total
unordered array	1	V	1	V^2
binary heap	$\log V$	$\log V$	$\log V$	$E \log V$
d-way heap	$\log_d V$	$d \log_d V$	$\log_d V$	$E \log_{E/V} V$
Fibonacci heap	1^\dagger	$\log V^\dagger$	1^\dagger	$E + V \log V$

† amortized

Bottom line.

- Array implementation optimal for dense graphs.
- Binary heap much faster for sparse graphs.
- 4-way heap worth the trouble in performance-critical situations.
-



<http://algs4.cs.princeton.edu>

4.4 SHORTEST PATHS

- ▶ *APIs*
- ▶ *shortest-paths properties*
- ▶ *Dijkstra's algorithm*
- ▶ ***edge-weighted DAGs***
- ▶ *negative weights*

Shortest paths in edge-weighted DAGs

```
public class AcyclicSP
{
    private DirectedEdge[] edgeTo;
    private double[] distTo;

    public AcyclicSP(EdgeWeightedDigraph G, int s)
    {
        edgeTo = new DirectedEdge[G.V()];
        distTo = new double[G.V()];

        for (int v = 0; v < G.V(); v++)
            distTo[v] = Double.POSITIVE_INFINITY;
        distTo[s] = 0.0;


        Topological topological = new Topological(G);
        for (int v : topological.order())
            for (DirectedEdge e : G.adj(v))
                relax(e);
    }
}
```

← topological order

Longest paths in edge-weighted DAGs

Formulate as a shortest paths problem in edge-weighted DAGs.

- Neglizho peshat.
- Gjej shortest paths.
- Neglizho peshat në rezultat.




equivalent: reverse sense of equality in `relax()`

longest paths input

5->4	0.35
4->7	0.37
5->7	0.28
5->1	0.32
4->0	0.38
0->2	0.26
3->7	0.39
1->3	0.29
7->2	0.34
6->2	0.40
3->6	0.52
6->0	0.58
6->4	0.93

shortest paths input

5->4	-0.35
4->7	-0.37
5->7	-0.28
5->1	-0.32
4->0	-0.38
0->2	-0.26
3->7	-0.39
1->3	-0.29
7->2	-0.34
6->2	-0.40
3->6	-0.52
6->0	-0.58
6->4	-0.93




Key
point.

Topological sort algorithm works even with negative weights.

Longest paths in edge-weighted DAGs: aplikimi

Parallel job scheduling. Le të jetë një bashkësi e punëve me kohëgjatje dhe precedence constraints, schedule punët (duke gjetur fillimin për secilën prej tyre) ashtu që të arrihet koha minimale e përfundimit, duke i respektuar constraints.

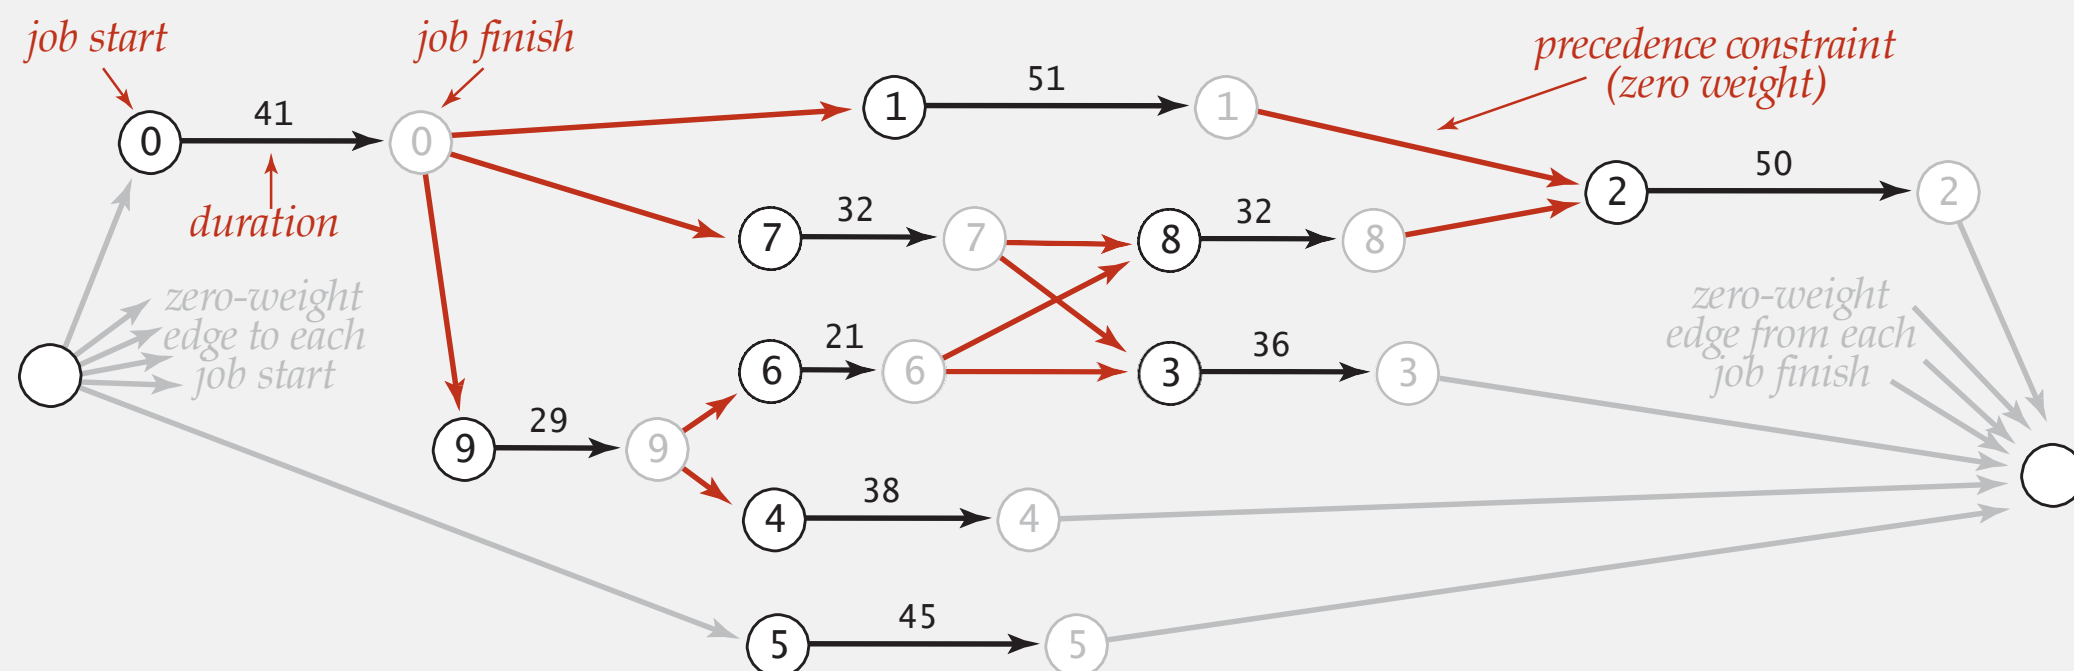


Critical path method

CPM. Për të zgjidhur parallel job-scheduling problem, krijo edge-weighted DAG:


- Kulmi burimor dhe përfundimtar.
- Dy kulme (fillimi dhe fundi) për çdo punë.
- Tri segmente për çdo punë.
 - fillimi - fundi (peshuar me kohëzgjatje)
 - burimi - fillimi (0 peshë)
 - fundi - përfundimi (0 peshë)
- Një segment për çdoprecedence constraint (0 peshë).

job	duration	must complete before		
0	41.0	1	7	9
1	51.0	2		
2	50.0			
3	36.0			
4	38.0			
5	45.0			
6	21.0	3	8	
7	32.0	3	8	
8	32.0	2		
9	29.0	4	6	

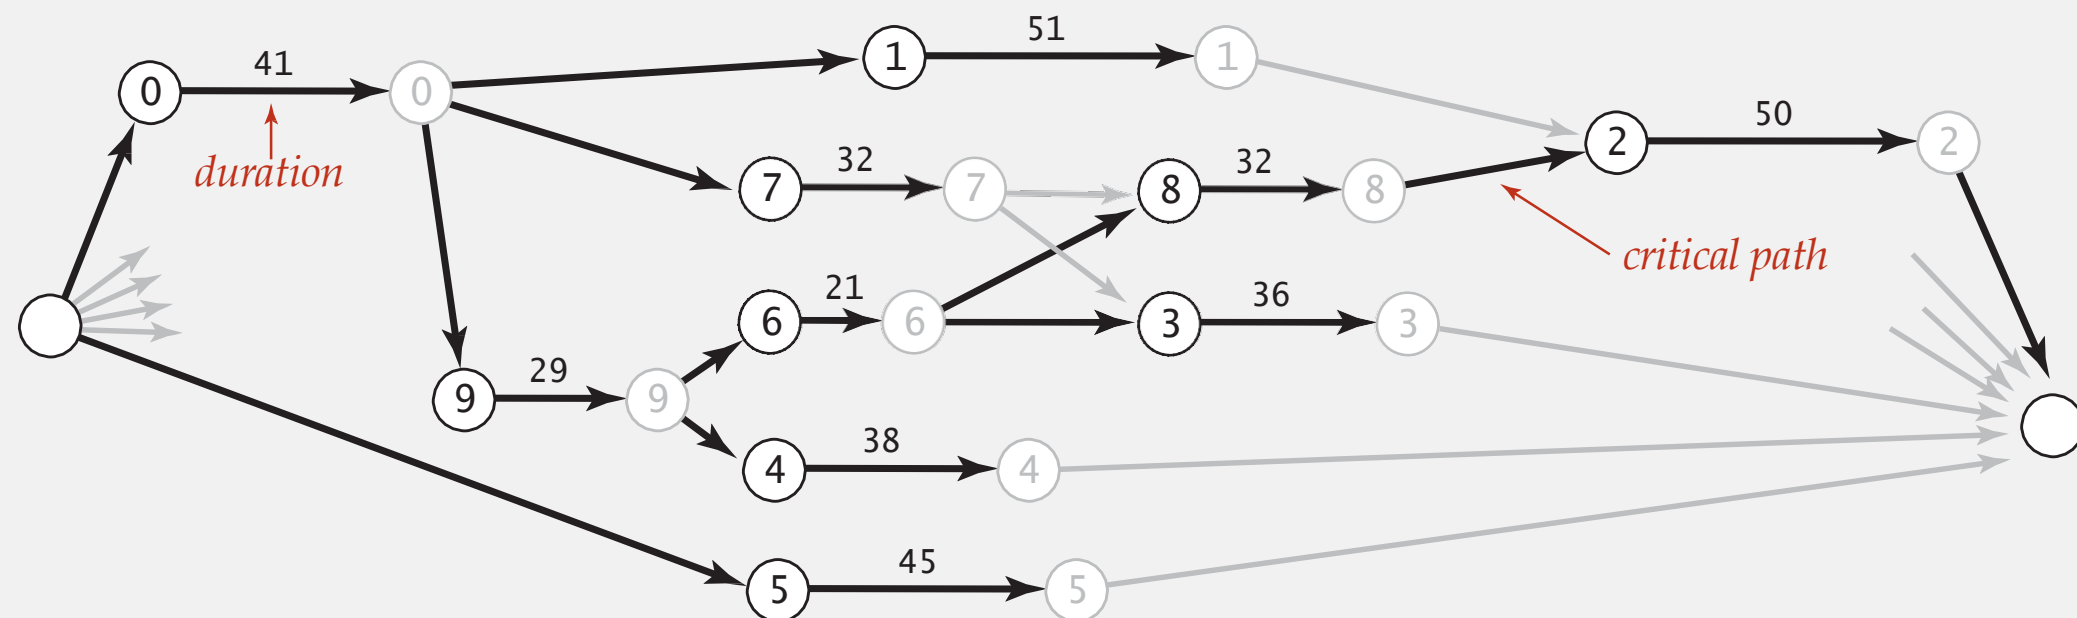


Critical path method

CPM. Use **longest path** from the source to schedule each job.



Parallel job scheduling solution





<http://algs4.cs.princeton.edu>

4.4 SHORTEST PATHS

▸ *APIs*

▸ *shortest-paths
properties*


▸ *Dijkstra's algorithm*

▸ *edge-weighted DAGs*

▸ ***negative weights***


Shortest paths with negative weights: failed attempts

Dijkstra. Nuk funksionon me pesha negative të segmenteve.



Dijkstra pëzgjedh kulmin 3 menjëherë pas 0. Por, shortest path prej 0 tek 3 është $0 \rightarrow 1 \rightarrow 2 \rightarrow 3$.

Re-peshimi. Shtimi i një konstante secilit segment nuk funksionon.



Shtimi i peshës për 8 tek të gjitha segmentet ndërron shortest path prej $0 \rightarrow 1 \rightarrow 2 \rightarrow 3$ në $0 \rightarrow 3$.


Përfundim. Duhet një algoritëm tjetër.

Negative cycles

Def. Një **negative cycle** është një directed cycle që ka shumën e peshës së segmenteve negative

digraph

4→5	0.35
5→4	-0.66
4→7	0.37
5→7	0.28
7→5	0.28
5→1	0.32
0→4	0.38
0→2	0.26
7→3	0.39
1→3	0.29
2→7	0.34
6→2	0.40
3→6	0.52
6→0	0.58
6→4	0.93



negative cycle $(-0.66 + 0.37 + 0.28)$

5→4→7→5

shortest path from 0 to 6

0→4→7→5→4→7→5...→1→3→6

Proposition. A SPT exists iff no negative cycles.

assuming all vertices reachable from s

Bellman-Ford algorithm

Bellman-Ford algorithm

Initialize $\text{distTo}[s] = 0$ and $\text{distTo}[v] = \infty$ for all other vertices.

Repeat V times:


- Relax each edge.**
-

```
for (int i = 0; i < G.V(); i++)  
    for (int v = 0; v < G.V(); v++)  
        for (DirectedEdge e : G.adj(v))  
            relax(e);
```

———— pass i (relax each edge)

Bellman-Ford algorithm demo

Repeat V times: relax all E edges.




0→1	5.0
0→4	9.0
0→7	8.0
1→2	12.0
1→3	15.0
1→7	4.0
2→3	3.0
2→6	11.0
3→6	9.0
4→5	4.0
4→6	20.0
4→7	5.0
5→2	1.0
5→6	13.0
7→5	6.0
7→2	7.0

an edge-weighted digraph

Bellman-Ford algorithm demo

Repeat V times: relax all E edges.




v	distTo[]	edgeTo[]
0	0.0	-
1	5.0	0→1
2	14.0	5→2
3	17.0	2→3
4	9.0	0→4
5	13.0	4→5
6	25.0	2→6
7	8.0	0→7


shortest t-paths tree from vertex s

Bellman-Ford algorithm: visualization

passes
4



7



10



13



SPT



Bellman-Ford algorithm: analysis

Bellman-Ford algorithm

Initialize $\text{distTo}[s] = 0$ and $\text{distTo}[v] = \infty$ for all other vertices.

Repeat V times:


- Relax each edge.**
-

Proposition. Dynamic programming algorithm computes SPT in any edge-weighted digraph with no negative cycles in time proportional to $E \times V$.

Pf idea. After pass i , found path that is at least as short as any shortest path containing i (or fewer) edges.

Bellman-Ford algorithm: practical improvement

Observation. If $\text{distTo}[v]$ does not change during pass i , no need to relax any edge pointing from v in pass $i+1$.

FIFO implementation. Maintain **queue** of vertices whose $\text{distTo}[]$ changed. 
be careful to keep at most one copy
of each vertex on queue (why?)

Overall effect.

- The running time is still proportional to $E \times V$ in worst case.
- But much faster than that in practice.

Single source shortest-paths implementation: cost summary

algorithm	restriction	typical case	worst case	extra space
topological sort	no directed cycles	$E + V$	$E + V$	V
Dijkstra (binary heap)	no negative weights	$E \log V$	$E \log V$	V
Bellman-Ford	no negative cycles	$E V$	$E V$	V
Bellman-Ford (queue-based)		$E + V$	$E V$	V

Remark 1. Directed cycles make the problem harder.

Remark 2. Negative weights make the problem harder.

Remark 3. Negative cycles makes the problem intractable.


Finding a negative cycle

Negative cycle. Add two methods to the API for SP.

boolean	hasNegativeCycle()	<i>is there a negative cycle?</i>
Iterable <DirectedEdge>	negativeCycle()	<i>negative cycle reachable from s</i>

digraph

4->5	0.35
5->4	-0.66
4->7	0.37
5->7	0.28
7->5	0.28
5->1	0.32
0->4	0.38
0->2	0.26
7->3	0.39
1->3	0.29
2->7	0.34
6->2	0.40
3->6	0.52
6->0	0.58
6->4	0.93




negative cycle (-0.66 + 0.37 + 0.28)

5->4->7->5

Finding a negative cycle

Observation. If there is a negative cycle, Bellman-Ford gets stuck in loop, updating `distTo[]` and `edgeTo[]` entries of vertices in the cycle.



Proposition. If any vertex v is updated in pass v , there exists a negative cycle (and can trace back `edgeTo[v]` entries to find it).

In practice. Check for negative cycles more frequently.